

SAKI SS 2021 Homework 1

Author: Andreas Thomas fo91juzi

Repository: <https://github.com/chronark/saki>

Summary

Classifying a small dataset of banking records into 7 categories using scikit-learn's Multinomial Naive-Bayes Classifier.

The dataset contained many unusable features so I focused on only 3 of them: "Buchungstext", "Verwendungszweck" and "Beguenstigter/Zahlungspflichtiger". I have done a little bit of cleaning by removing punctuation and transforming everything to lower case. This did not make a significant difference though. The relevant features were concatenated into one and tokenized using a CountVectorizer.

The other features were either very similar across all records or did not carry any meaning for this classification. I tried using the dates as well thinking that recurring transactions for salary or rent could be detected but it did not make a significant difference because the dataset is rather small.

Evaluation

Using the count vectorizer and a multinomial naive bayes algorithm I achieved ~ 88% accuracy. I used a multinomial classifier because multinomial is well suited and the input variables are discrete numbers, see: https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html. I have chosen not to use a tfidf transformer after testing it [Cell 7]. It is generally not suited for this dataset because the type of used features is rather short and most tokens are unique [Cell 5].

As you can see in the confusion matrix [Cell 9] most incorrect classifications are in the categories "leisure" and "standardOfLiving".

Notebook

notebook

May 11, 2021

```
[1]: import pandas as pd
from sklearn.naive_bayes import MultinomialNB
from tqdm import tqdm
from sklearn.metrics import accuracy_score, \
    →classification_report, plot_confusion_matrix
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.pipeline import Pipeline
import numpy as np
from sklearn.model_selection import cross_val_score
```

```
[2]: # Represents the proportion of the dataset used for training. Remaining will be \
    →used for validation.
TRAIN_SIZE = 0.7
```

```
[3]: # Load dataset

source = pd.read_csv("./SAKI Exercise 1 - Transaction Classification - Data Set.
    →csv", delimiter=";")
source.drop("Unnamed: 0", axis=1, inplace=True)

source.head()
```

```
[3]:   Auftragskonto Buchungstag Valutadatum      Buchungstext \
0      89990201.0  28.07.2016  28.07.2016      Lohn / Gehalt
1      89990201.0  27.07.2016  27.07.2016           Miete
2      89990201.0  21.07.2016  21.07.2016          Bargeld
3      89990201.0  20.07.2016  20.07.2016  Lebensmittel / Getraenke
4      89990201.0  18.07.2016  18.07.2016      Spontanausgabe

                                Verwendungszweck \
0  Gehalt Adorsys GmbH & Co. KG End-To-End-Ref.: ...
1  Byladem1Sbt De12773501123456789889 Miete Beuth...
2  21.07/16.34Uhr Nuernberg All Eur 70,00 Geb.Eur...
3  2831 Edeka Neubauer Nuernb.//Nuernb 2016-07-20...
4                                     Amazon
```

	Beguenstigter/Zahlungspflichtiger	Kontonummer	BLZ \
0	Adorsys GmbH & Co. KG	7807800780	25190001
1	Georg Tasche	DE31251900019123456780	VOHADE2HXXX
2	Bargeld	9999900780	25190001
3	Kartenzahlung	9736000780	25190001
4	neue Playstation	9988776655	25125100

	Betrag	Waehrung	label
0	2000.00	EUR	income
1	-670.00	EUR	living
2	-70.00	EUR	private
3	-73.21	EUR	standardOfLiving
4	-363	EUR	leisure

```
[4]: # Extract relevant features and concatenate them to a single lower case string
      ↳per feature

# Removing some characters could be helpful.
# However these are almost always part of a word and are not surrounded by
      ↳whitespace.
# Hence they are not detected as a single token by the count tokenizer and it
# increases the accuracy by a very tiny amount, though that might change with a
# bigger dataset
def remove_stopwords(feature: str)-> str:
    feature = feature.replace(".", "")
    feature = feature.replace(",", "")
    feature = feature.replace(":", "")
    feature = feature.replace("/", "")
    feature = feature.replace("&", "")
    return feature

# Concatenate all relevant features into one big string
features = []
for i in range(len(source.index)):
    combined = []
    combined.append(source.loc[i, "Buchungstext"])
    combined.append(source.loc[i, "Verwendungszweck"])
    combined.append(source.loc[i, "Beguenstigter/Zahlungspflichtiger"])

    feature = remove_stopwords(" ".join(combined))

    features.append(feature.lower())

# Extract labels as well
labels = list(source["label"])
```

```
print("Example of features as a single string:")
print(features[0])
```

Example of features as a single string:

```
lohn gehalt gehalt adorsys gmbh co kg end-to-end-ref notprovided
kundenreferenz nsct16033000136600000000000000000001 gutschrift adorsys gmbh co
kg
```

```
[5]: # I would like to know how many tokens each feature has to judge whether or not,
      ↪ using a tfidf transformer makes sense
def inspect_tokens(features: [str]) -> None:
    tokens_per_feature = []
    unique_tokens_per_feature = []
    for feature in features:
        tokens = feature.split(" ")
        tokens_per_feature.append(len(tokens))
        unique_tokens_per_feature.append(len(set(tokens)))

    print("Maximum tokens:", max(tokens_per_feature))
    print("Average tokens:", np.average(tokens_per_feature))
    print("Average unique tokens", np.average(unique_tokens_per_feature) )

inspect_tokens(features)
```

Maximum tokens: 19

Average tokens: 11.311004784688995

Average unique tokens 10.311004784688995

We can see that there really aren't that many duplicate tokens in each feature, so I think a tfidf transformer will not help us a lot. I am testing this directly further down.

```
[6]: # Helper function to create a new pipeline with or without tfidf transformer
def create_pipeline(with_tfidf: bool) -> Pipeline:
    steps = [ ("vectorizer", CountVectorizer())]
    if (with_tfidf is True):
        steps.append(("tfidf", TfidfTransformer()))
    steps.append(("nb", MultinomialNB()))
    return Pipeline(steps)

# Split our dataset
def split_data(features: [str], labels: [str]):
    return train_test_split(features, labels, train_size=TRAIN_SIZE,
        ↪ shuffle=True)
```

```
[7]: # Runs a bunch of predictions with random train/test splits to get an average
      ↪ accuracy.
```

```

def do_random_runs(with_tfidf: bool) -> [float]:
    accuracies = []
    for i in tqdm(range(1000)):
        x_train, x_test, y_train, y_test = split_data(features, labels)
        classifier = create_pipeline(with_tfidf=with_tfidf)
        classifier.fit(x_train, y_train)

        accuracies.append(accuracy_score(y_test, classifier.predict(x_test)))
    return accuracies

# Plot accuracies inside a subplot
def plot(ax, accuracies: [float]) -> None:
    accuracies.sort(reverse=True)
    mean = np.mean(accuracies)
    std = np.array(accuracies).std()

    ax.set_xlabel("mean={:.3f}, std={:.3f}".format(mean, std))
    x = range(0, len(accuracies))

    y = accuracies

    ax.scatter(x=x, y=y, label="Accuracy")
    ax.plot(x, [mean]*len(x), label="Mean", linestyle="--")

# Calculate and plot the difference when using a tfidf transformer or not
def compare_tfidf(features: [str], labels: [str]) -> None:
    fig, (ax1, ax2) = plt.subplots(1, 2, sharex=True, sharey=True)
    fig.suptitle("[Figure 1] Comparison with and without tfidf transformer over 1000 random predictions")

    # 1. with tfidf
    plot(ax1, do_random_runs(True))
    ax1.set_title("with tfidf")
    ax1.set_ylabel("Accuracy")

    # 2. without tfidf
    plot(ax2, do_random_runs(False))
    ax2.set_title("without tfidf")

    plt.show()

compare_tfidf(features, labels)


```

```

100%|      | 1000/1000 [00:04<00:00, 202.39it/s]
100%|      | 1000/1000 [00:03<00:00, 271.04it/s]

```

[Figure 1] Comparison with and without tfidf transformer over 1000 random predictions



Initially I thought using a tfidf transformer would not increase our accuracy a lot because our features have fewer than 20 tokens that do not repeat often. But it actually decreases the accuracy a little bit.

A mean accuracy of ~ 0.88 and standard deviation of ~ 0.04 is acceptable however the big range of accuracies depending on the shuffled train/test split is a result of the very small dataset.


[8]: *# Final model*

```
x_train, x_test, y_train, y_test = split_data(features, labels)
classifier = create_pipeline(with_tfidf=False)
classifier.fit(x_train, y_train)
```

[8]: Pipeline(steps=[('vectorizer', CountVectorizer()), ('nb', MultinomialNB())])

[9]: *# Let's look at the confusion matrix for our model*

```
plot = plot_confusion_matrix(classifier, x_test, y_test, cmap=plt.cm.Blues)
plot.ax_.set_title("[Figure 2] Confusion matrix")
plt.show()
```



```
[10]: # And print some statistics
print(classification_report(y_test, classifier.predict(x_test)))
```

	precision	recall	f1-score	support
finance	1.00	0.93	0.96	14
income	1.00	1.00	1.00	7
leisure	0.77	1.00	0.87	10
living	1.00	0.67	0.80	9
private	1.00	1.00	1.00	5
standardOfLiving	0.89	0.94	0.92	18
accuracy			0.92	63
macro avg	0.94	0.92	0.93	63
weighted avg	0.93	0.92	0.92	63